Superpowered Optimization in Python With Gurobi and Anaconda

Renan Garcia, Ph.D., Gurobi Optimization
Michael Grant, Ph.D., Continuum Analytics

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM® ANALYTICS

GUROBI OPTIMIZATION

Renan Garcia, Ph.D.

- Optimization Support Engineer at Gurobi Optimization
- Ph.D. in Industrial and Systems Engineering, Georgia Tech
- Expert in optimization modeling and software development
- Over a decade of experience implementing decision support systems

CONTINUUM® ANALYTICS

GUROBI OPTIMIZATION

Michael C. Grant, Ph.D.

- Senior Solution Architect at Continuum Analytics
- Ph.D. in Electrical Engineering, Stanford University
- Subject matter expert in nonlinear convex optimization
- Developer of the CVX modeling framework

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Overview

- The Proposition
- Introducing Gurobi Optimization
- About Continuum Analytics & Anaconda
- The Gurobi Python module
- A simple MIP model using Gurobi & Spyder
- Portfolio optimization using Gurobi, Jupyter, Pandas, & Bokeh
- Wrap up / Call to action

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

The Proposition

- The Gurobi Python module provides the intuitive feel of a *modeling* language inside a powerful *programming* language
- The Anaconda platform delivers a full application development stack as well as a great interactive environment for data analysis and modeling
- Bringing developers and analysts/modelers together using Python means collaboration has never been easier

The Gurobi Optimizer

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Inside the Gurobi Optimizer

- Algorithms for continuous optimization (LP, QP, QCP)
  - Simplex, Barrier
- Algorithms for discrete optimization (MIP, MIQP, MIQCP)
  - Parallel branch-and-bound
- Other algorithms
  - Presolve, parameter tuning, irreducible inconsistent subsystems…
- Command-line interface, full-featured interactive shell
- Programming interfaces
  - C, C++, Java, .NET, Python, MATLAB, R

Gurobi Model Components

- Decision variables
- Objective function
  - minimize $\mathbf{x}^T\mathbf{Q}\mathbf{x} + \mathbf{c}^T\mathbf{x} + \alpha$
- Constraints
  - $\mathbf{A}\mathbf{x} = \mathbf{b}$ (linear constraints)
  - $\mathbf{l} \leq \mathbf{x} \leq \mathbf{u}$ (bound constraints)
  - some x_h integral (integrality constraints)
  - some x_i lie within second order cones (cone constraints)
  - $\mathbf{x}^T\mathbf{Q}_j\mathbf{x} + \mathbf{q}_j^T\mathbf{x} \leq \beta_j$ (quadratic constraints)
  - some x_k in SOS (special ordered set constraints)

Example: Mixed Integer Program

- Decision variables
- Objective function
  - minimize $\mathbf{x}^T\mathbf{Q}\mathbf{x} + \mathbf{c}^T\mathbf{x} + \alpha$
- Constraints
  - $\mathbf{A}\mathbf{x} = \mathbf{b}$ (linear constraints)
  - $\mathbf{l} \leq \mathbf{x} \leq \mathbf{u}$ (bound constraints)
  - some x_h integral (integrality constraints)
  - some x_i lie within second order cones (cone constraints)
  - $\mathbf{x}^T\mathbf{Q}_j\mathbf{x} + \mathbf{q}_j^T\mathbf{x} \leq \beta_j$ (quadratic constraints)
  - some x_k in SOS (special ordered set constraints)

Benchmarks – Open Source MIP

- Mittelmann MIPLIB2010 tests, P=1 (>1X means Gurobi wins):
  - http://plato.la.asu.edu/bench.html

	January 2012	January 2013	January 2015	November 2015	% Solved
CBC	10X	13X	17X	26X	55%
SCIP	6X	7X	7X	9X	80%
GLPK	22X	27X	-	-	1%
LPSOLVE	19X	24X	-	-	6%

- GLPK and LPSOLVE are not currently tested in the public benchmarks

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Why Gurobi?

- Performance leader
  - Average 2X improvement with each major release
- Consistent track record of innovation
  - First cloud offering, client-server support, distributed algorithms, …
- Outstanding support
  - Direct access to Ph.D.-level optimization experts
- No surprises
  - Flexible licensing and transparent pricing
- Strong academic offering
  - Free full-featured academic licenses, Take Gurobi with You program

About Continuum Analytics & Anaconda

GUROBI
OPTIMIZATION

ANACONDA®

Continuum Analytics

- Founded in 2012 by Travis Oliphant and Peter Wang
- Key contributors to the Python open source ecosystem
  - Travis: primary developer, NumPy; founding contributor, SciPy
  - PyTables, Pandas, Jupyter/IPython, Matplotib
  - New projects: Blaze, Numba, Conda
- Commercial efforts:
  - Enterprise software & support
  - Training: Python for Science, Python for Finance
  - Consulting: custom software development

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Anaconda for Data Science

Empowering Everyone on the Team

Data Scientist

- Advanced analytics with Python & R
- Simplified library management
- Easily share data science notebooks & packages

Developer

- Support for common APIs & data formats
- Common language with data scientists
- Python extensibility with C, C++, etc.

Ops

- Validated source of up-to-date packages including indemnification
- Agile Enterprise Package Management
- Supported across platforms

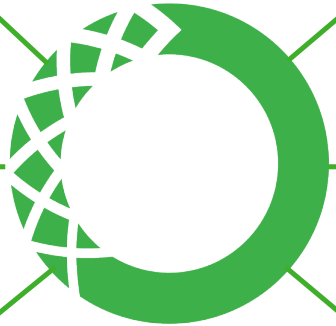

Data Engineer

- Powerful & efficient libraries for data transformations
- Robust processing for noisy dirty data
- Support for common APIs & data formats

Business Analyst

- Collaborative interactive analytics with notebooks
- Rich browser based visualizations
- Powerful MS Excel integration

Computational Scientist

- Rich set of advanced analytics
- Trusted & production ready libraries for numerics
- Simplified scale up & scale out on clusters & GPUs

Anaconda for Data Science

Empowering Everyone on the Team

Data Scientist

- Advanced analytics with Python & R
- Simplified library management
- Easily share data science notebooks & packages

Developer

- Support for common APIs & data formats
- Common language with data scientists
- Python extensibility with C, C++, etc.

Ops

- Validated source of up-to-date packages including indemnification
- Agile Enterprise Package Management
- Supported across platforms

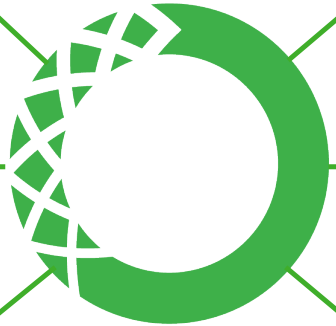

Data Engineer

- Powerful & efficient libraries for data transformations
- Robust processing for noisy dirty data
- Support for common APIs & data formats

Business Analyst

- Collaborative interactive analytics with notebooks
- Rich browser based visualizations
- Powerful MS Excel integration

Computational Scientist

- Rich set of advanced analytics
- Trusted & production ready libraries for numerics
- Simplified scale up & scale out on clusters & GPUs

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Why Python?

- Designed for readability
- Functional, object-oriented, declarative
- Great glue language for legacy C/Fortran
- Rich standard library
- Incredible community library
- Great online resources
- Great name

"The second-best language for everything"
— Peter Wang, CTO

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

The Anaconda Python Distribution

- Windows, Mac, Linux; 32/64-bit
- Python 2.x, 3.x — your choice (though do pick 3.x!)
- ~150 packages in the default installation
- **`conda`** package and environment management
  - ~350 packages available with a simple **`conda install`**
  - Create isolated “silos” for multiple projects
- Like R? Anaconda has that too

ANACONDA®

conda package management

Install	`conda install pandas numpy=1.9 gurobi`
Update	`conda update gurobi`
List	`conda list`
Search	`conda search gurobi`
Remove	`conda remove matplotlib`

conda environment management

Create	`conda create -n lp_project anaconda gurobi`	
Select	`source activate lp_project` `activate lp_project`	*(Linux/Mac)* *(Windows)*
Return to default	`source deactivate` `deactivate`	*(Linux/Mac)* *(Windows)*
Remove	`conda remove -n lp_project --all`	

Obtaining Anaconda

From Continuum: https://www.continuum.io/downloads

From Gurobi: http://www.gurobi.com/downloads/get-anaconda

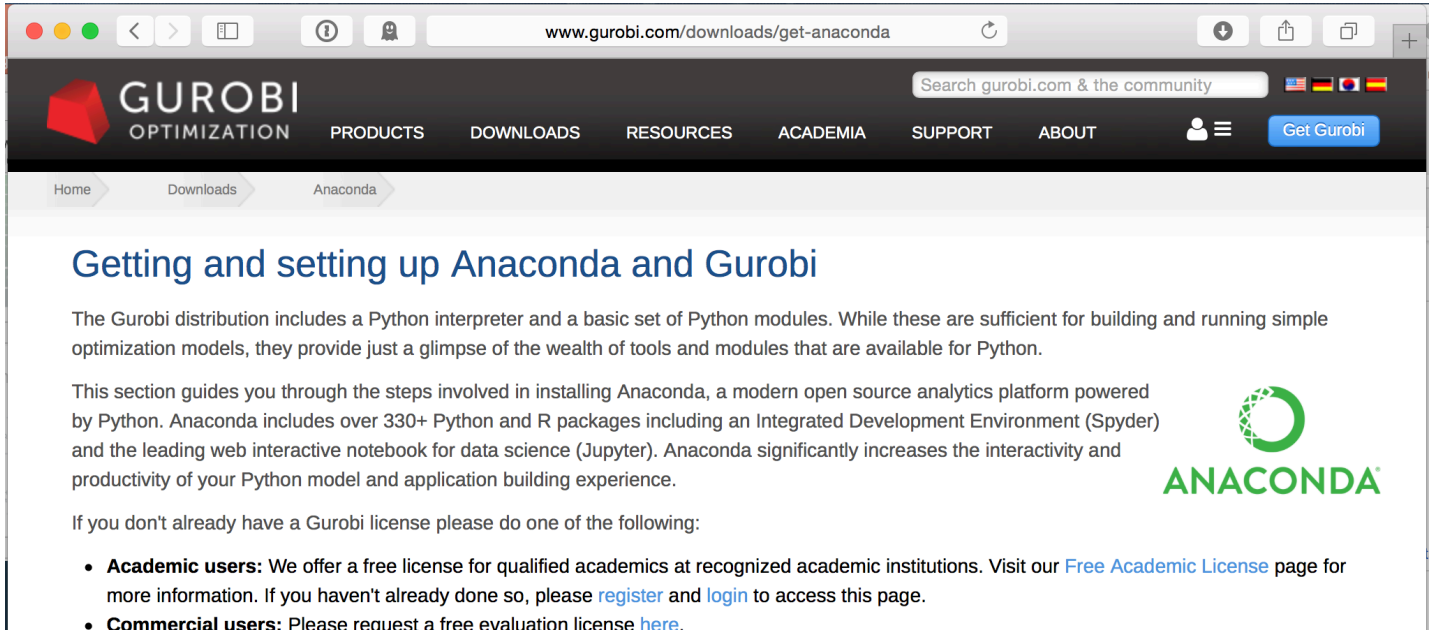

The Gurobi Python Module

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Building Models with Gurobi

- Option 1: use a modeling language
  - Very easy to build models
    - Optimization modeling constructs built into language
  - Attractive choice for non-programmers
- Option 2: use a full programming language
  - Much more powerful and flexible development environment
    - Complete access to solver functionality
    - Richer set of language features
  - Natural choice when deploying and/or integrating with applications

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Gurobi Python Environment

- High-level optimization modeling constructs embedded in Python
- Design goals:
  - Require minimal programming skills to get started
  - Bring "feel" of a modeling language to the Python interface
  - Allow for code that is easy to write and maintain
  - Maintain unified design across all of our interfaces
  - Remain lightweight and efficient compared to solver alone
  - Support all solver and programming needs

Essential Gurobi Python Constructs

- Objects represent model components (`Model`, `Var`, `Constr`)
- Overloaded operators
  - Arithmetic (+, -, ×, ÷), constraints (≤, =, ≥)
- Aggregate sum operator (`quicksum`)
- Python provides the rest for representing data, indices and subscripts
  - Lists, tuples, dictionaries, loops, generator expressions, …

Ex: $x_i + y_i \leq 5, \forall i \in I$ $\Leftrightarrow$

```
for i in I:
    m.addConstr(x[i] + y[i] <= 5)
```

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Building a Gurobi Model: Six Steps

Create a model	`m = Model()`
Add variables	`x = m.addVar(...)`
Commit changes	`m.update()`
Set objective	`m.setObjective(...)`
Add constraints	`m.addConstr(...)`
Optimize	`m.optimize()`

Installing the Gurobi Python Module

Add the channel	`conda config --add channels gurobi`
Install Gurobi	`conda install gurobi`
Retrieve license	`grbgetkey xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx`
Test license	`>>> import gurobipy`
Test token server	`gurobi_cl --tokens`

A Simple MIP Model using the Spyder IDE

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Spyder

- **S**cientific **Py**thon **D**evelopment **E**nvi**R**onment
- Familiar *integrated development environment* (IDE) for users of MATLAB, Visual Studio, etc.
- Built-in editor, file explorer, documentation browser
- Graphical interface for debugging (ipdb)
- IPython console for interactive code execution

http://pythonhosted.org/spyder

Spyder Demo

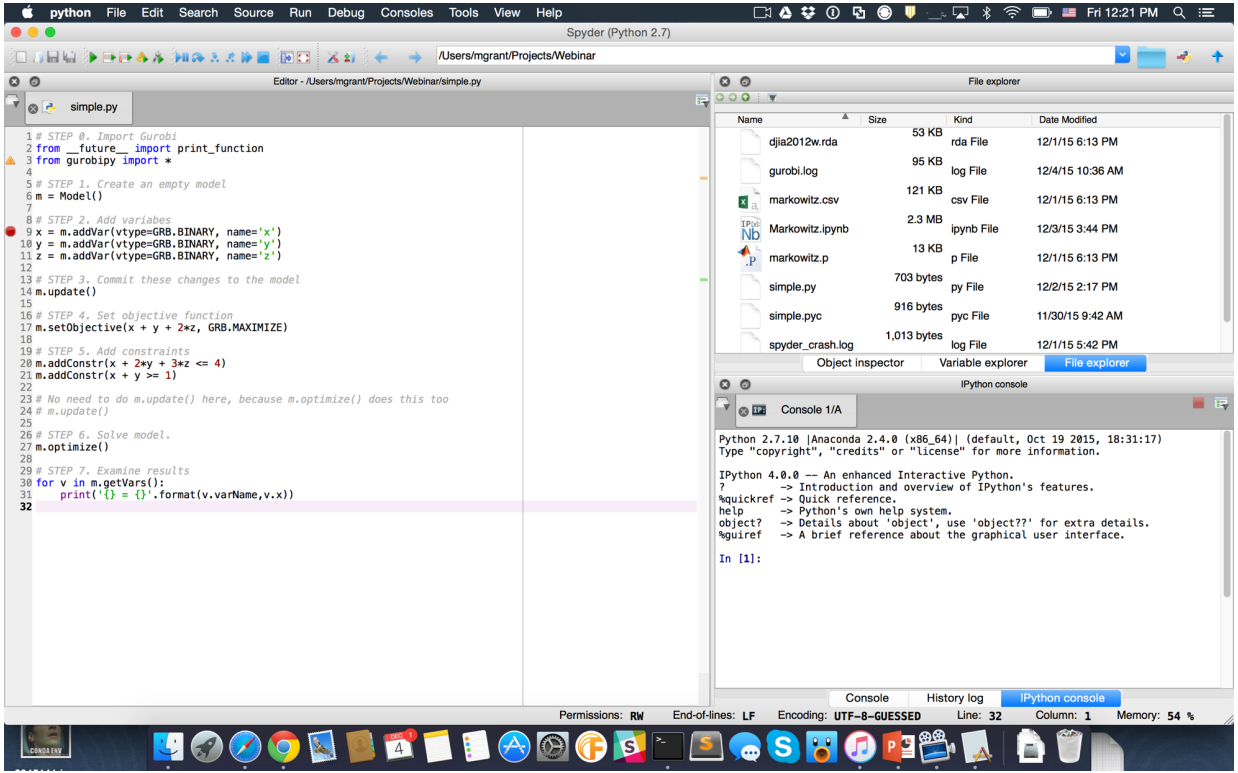
python File Edit Search Source Run Debug Consoles Tools View Help
Fri 12:21 PM
Spyder (Python 2.7)
/Users/mgrant/Projects/Webinar
Editor - /Users/mgrant/Projects/Webinar/simple.py
simple.py
# STEP 0. Import Gurobi
from __future__ import print_function
from gurobipy import *
# STEP 1. Create an empty model
m = Model()
# STEP 2. Add variabes
x = m.addVar(vtype=GRB.BINARY, name='x')
y = m.addVar(vtype=GRB.BINARY, name='y')
z = m.addVar(vtype=GRB.BINARY, name='z')
# STEP 3. Commit these changes to the model
m.update()
# STEP 4. Set objective function
m.setObjective(x + y + 2*z, GRB.MAXIMIZE)
# STEP 5. Add constraints
m.addConstr(x + 2*y + 3*z <= 4)
m.addConstr(x + y >= 1)
# No need to do m.update() here, because m.optimize() does this too
# m.update()
# STEP 6. Solve model.
m.optimize()
# STEP 7. Examine results
for v in m.getVars():
print('{} = {}'.format(v.varName,v.x))
File explorer
Name
Size
Kind
Date Modified
djia2012w.rda 53 KB rda File 12/1/15 6:13 PM
gurobi.log 95 KB log File 12/4/15 10:36 AM
markowitz.csv 121 KB csv File 12/1/15 6:13 PM
Markowitz.ipynb 2.3 MB ipynb File 12/3/15 3:44 PM
markowitz.p 13 KB p File 12/1/15 6:13 PM
simple.py 703 bytes py File 12/2/15 2:17 PM
simple.pyc 916 bytes pyc File 11/30/15 9:42 AM
spyder_crash.log 1,013 bytes log File 12/1/15 5:42 PM
Object inspector
Variable explorer
File explorer
IPython console
Console 1/A
Python 2.7.10 |Anaconda 2.4.0 (x86_64)| (default, Oct 19 2015, 18:31:17)
Type "copyright", "credits" or "license" for more information.
IPython 4.0.0 -- An enhanced Interactive Python.
? -> Introduction and overview of IPython's features.
%quickref -> Quick reference.
help -> Python's own help system.
object? -> Details about 'object', use 'object??' for extra details.
%guiref -> A brief reference about the graphical user interface.
In [1]:
Console
History log
IPython console
Permissions: RW
End-of-lines: LF
Encoding: UTF-8-GUESSED
Line: 32
Column: 1
Memory: 54 %

Portfolio Optimization in Jupyter

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Modern Portfolio Theory In One Slide

- First offered by Harry Markowitz in 1952
- Recognized by a Nobel Memorial Prize in Economics
- Associates with each stock an *expected return* and a risk or *volatility* measure, and a matrix of *correlations* between pairs of stocks
- The goal is to study the tradeoff between *risk* and *return*
  - Minimize risk for a given target (expected) return
  - Maximize return given an upper bound on risk
- The most basic models are *quadratic programs* (QPs)

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

(okay, two slides)

$$\begin{array}{ll} \text{minimize} & x^T \Sigma x \\ \text{subject to} & \vec{1}^T x = 1 \\ & \rho^T x = \gamma \\ & x \geq 0 \end{array} \qquad \begin{array}{ll} \text{minimize} & \sum_{i=1}^{n} \sum_{j=1}^{n} \sigma_{ij} x_i x_j \\ \text{subject to} & \sum_{i=1}^{n} x_i = 1 \\ & \sum_{i=1}^{n} \rho_i x_i = \gamma \\ & x_i \geq 0 \quad i = 1, 2, \ldots, n \end{array}$$

x_i	amount of investment allocated to stock i
ρ_i	expected return for stock i
σ_{ii}	variance of the performance for stock i
$\sigma_{ij} / \sqrt{\sigma_{ii}\sigma_{jj}}$	correlation between stocks i and j

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Jupyter (formerly IPython)

- A web application for interactive data analysis
- Create notebooks containing live code, text, equations, and visualizations
- A great way to prototype new algorithms, document experiments, share reproducible results
- Originally for Python, now there are kernels for over 50 different languages

http://www.jupyter.org

jupyter

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Pandas

- Fast and efficient DataFrame and Series objects
- Read/write CSV, HDF5, Excel, SQL, plain text
- Missing data handling
- Slicing, fancy indexing, subsetting
- Merges, joins
- Split-apply-combine operations (groupby)

http://pandas.pydata.org

pandas $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$

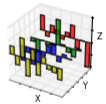

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Bokeh

Interactive visualization for web browsers

- High performance with streaming data and big data
- In-browser interactivity on otherwise *static* pages
- Bindings for Python, Scala, Julia, R

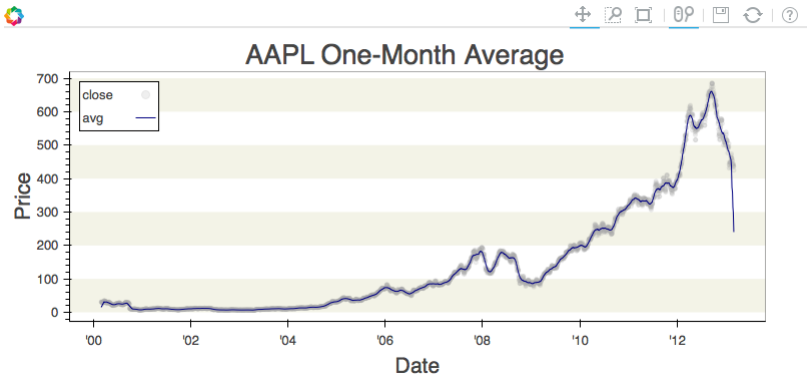

http://bokeh.pydata.org

Jupyter Demo

https://notebooks.anaconda.org/mcg/portfolio

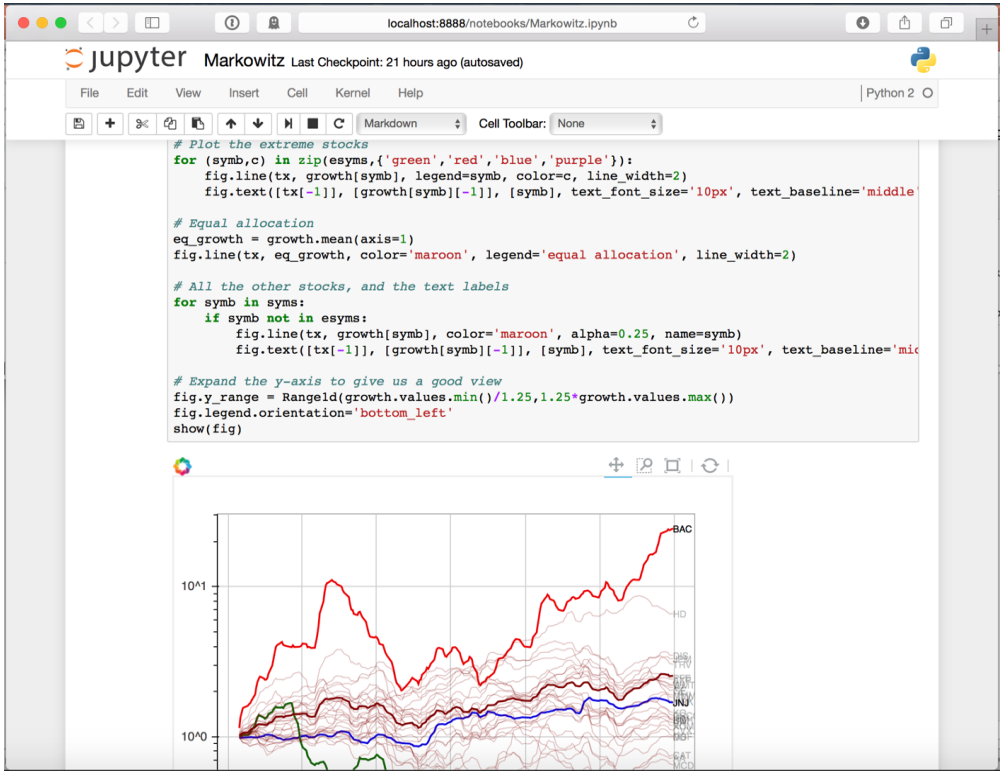

Closing thoughts

GUROBI
OPTIMIZATION

ANACONDA®

CONTINUUM ANALYTICS

GUROBI OPTIMIZATION

Developers vs. Analysts?

Analysts benefit from Jupyter:

- Interactive, iterative development
- Notebooks combine code, results, visualizations, documentation

Developers benefit from Spyder:

- Traditional IDE organization
- Graphical debugging

Both benefit from *straightforward collaboration:*

- Code from notebooks readily inserted into standard Python modules

Wrap up

- The Gurobi Python interface brings the expressive power of a modeling language into the powerful Python ecosystem
- Spyder offers a traditional IDE for Python development, including full-featured editing and graphical debugging
- IPython and the Jupyter notebook system provide a powerful paradigm for interactive development, collaboration, and sharing
- The Anaconda platform delivers a suite of powerful open-source tools and connects to world-class commercial software like Gurobi
- Need training, consulting, enterprise support? Continuum can help

Thank you!

Download Gurobi and/or Anaconda and try them for yourself:

www.gurobi.com/get-anaconda

Email: sales@gurobi.com

Twitter: @Gurobi

Sales: sales@continuum.io

Twitter: @ContinuumIO